



Chapter3

Understanding Big Data Technology Foundations

Lamda Architecture

Basanta Joshi, PhD

Asst. Prof., Dept of Electronics and Computer Engineering

Program Coordinator, MSc in Information and Communication Engineering

Member, Laboratory for ICT Research and Development (LICT)

Member, Research Management Cell (RMC)

Institute of Engineering

basanta@ioe.edu.np

<http://www.basantajoshi.com.np>

<https://scholar.google.com/citations?user=iocLiGcAAAAJ>

https://www.researchgate.net/profile/Basanta_Joshi2



The world is changing !

The model of Generating/Consuming Data has changed !.

Old Model: few companies are generating data, all others are consuming data



New Model: all of us are generating data, and all of us are consuming data

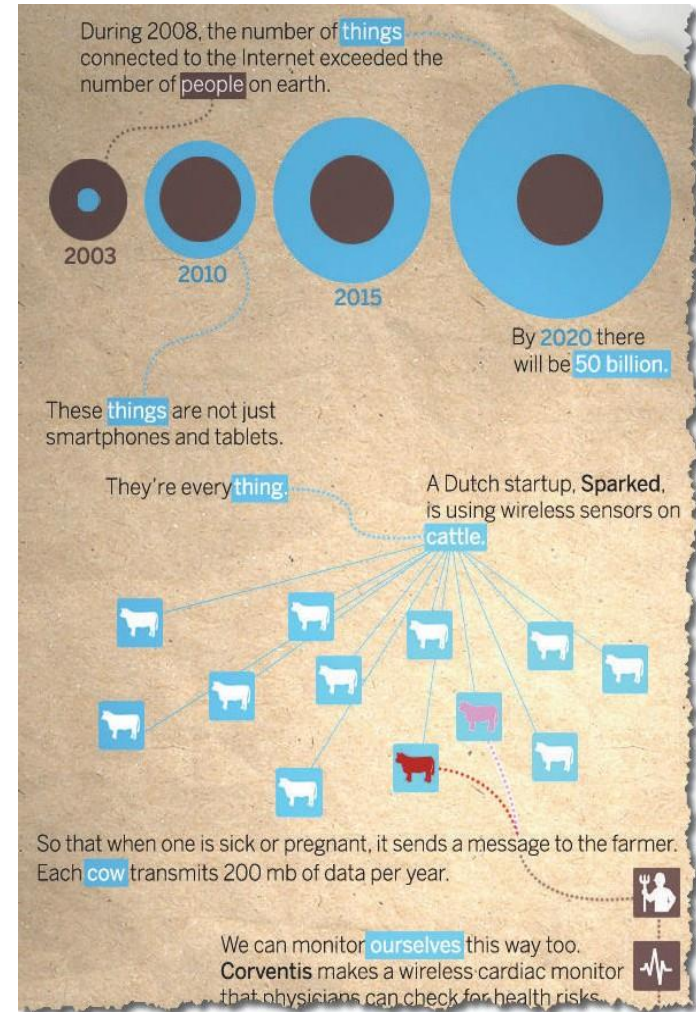





Internet Of Things – Sensors are/will be everywhere

There are more devices tapping
into the internet than people on
earth

How do we prepare our
systems/architecture for the
future?





2020 IoT Insights

- With 1.3 billion projected subscriptions by 2023, IoT is about to experience another boost by the 5G technology.
- By 2022, **Google Home will have the largest IoT devices market share**, at 48%.
- The average number of connected devices **per household in 2020 will be 50**.
- By 2021, **35 billion IoT devices** will be installed around the world.
- The number of connected devices **in 2020 is predicted to hit 50 billion**.
- <https://techjury.net/blog/how-many-iot-devices-are-there/#gref>



The world is changing

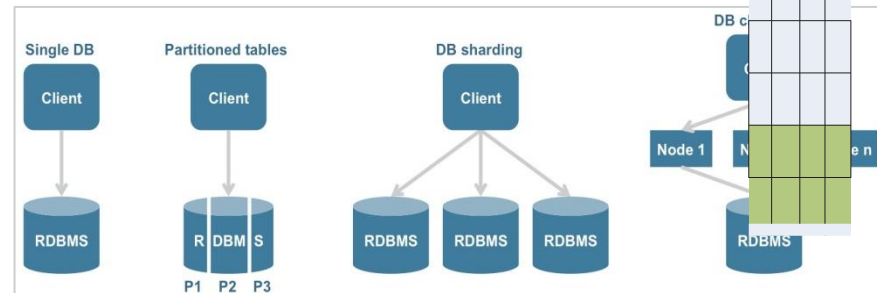
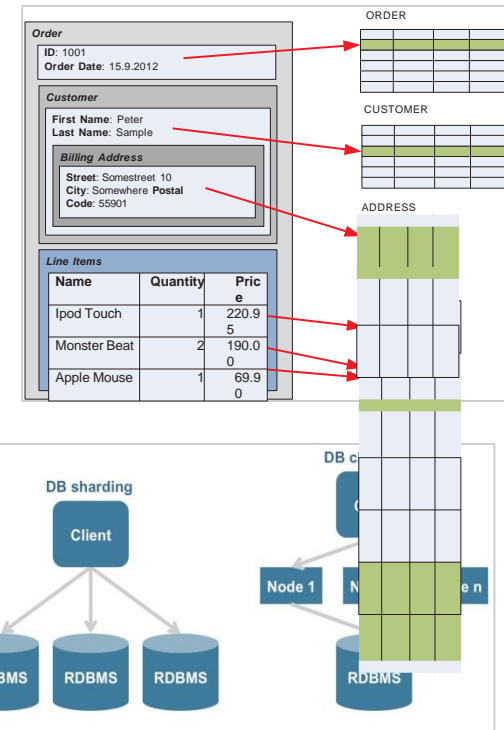
new data stores



Problem of traditional (R)DBMS

approach:
■ Complex object
graph

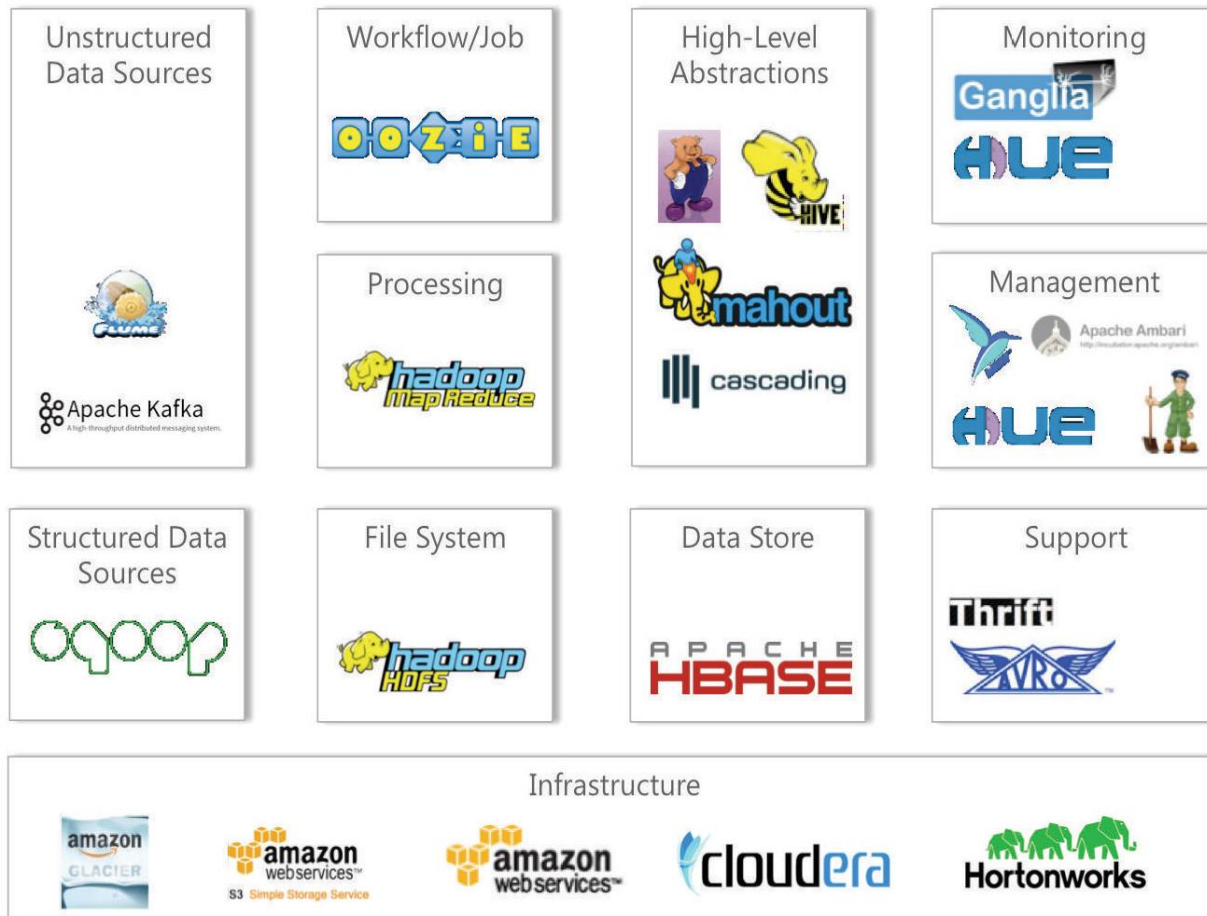
- Schema evolution
- Semi-structured data
- Scaling



Polyglot persistence

- Using multiple data storage technologies (RDBMS + NoSQL + NewSQL + In-Memory)

■ The world is changing ! New platforms evolving (i.e. Hadoop Ecosystem)



Data as an Asset – Store everything?



Data is just too valuable to delete!
We must store anything!

It depends ...
Big Data technologies allow **to store the raw information** from new and existing data sources so that you can **later use it** to create **new data-driven products**, which you haven't ~~thought about~~ **today!**



Nonsense!
Just store the data **you know you need today!**





AGENDA

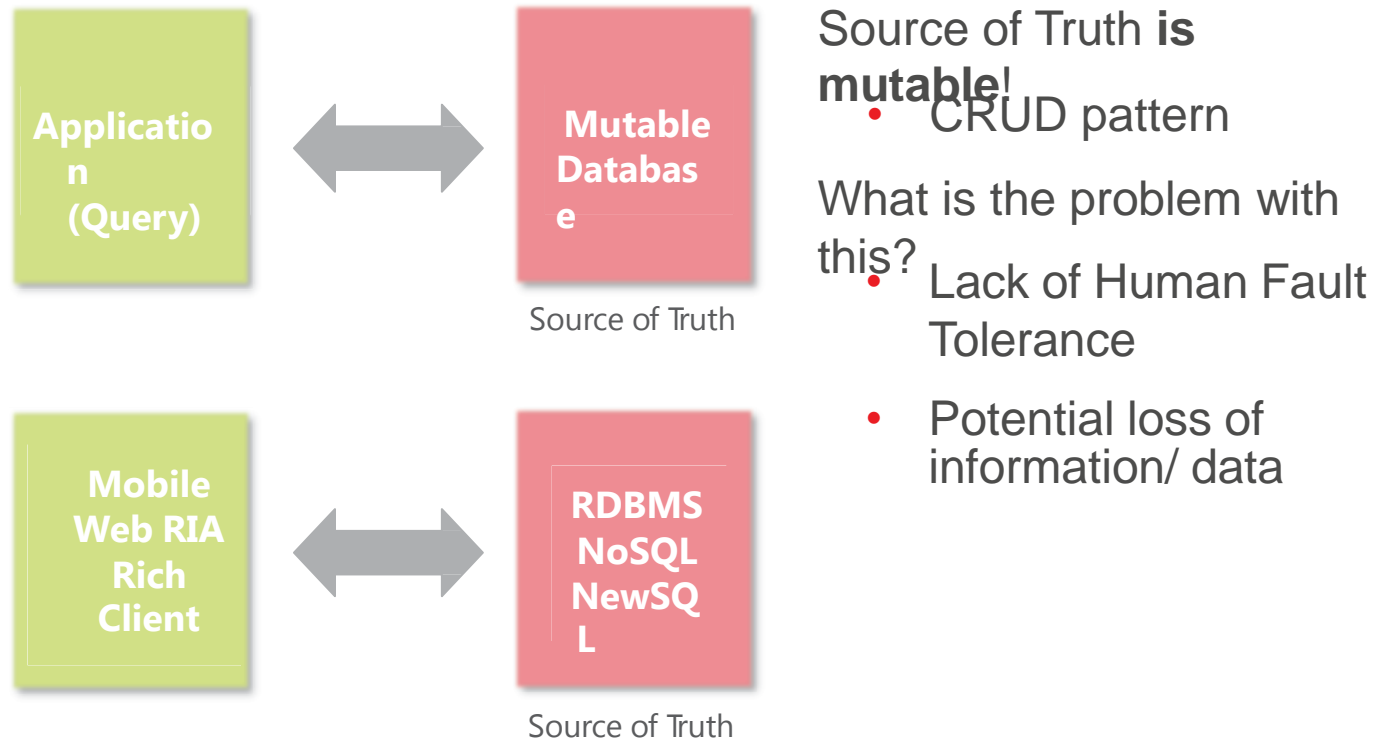
1. Big Data and Fast Data, what is it?
- 2. Architecting (Big) Data Systems**
3. The Lambda Architecture
4. Use Case and the Implementation
5. Summary and Outlook



What is a data system?

- A (data) system that manages the **storage and querying of data** with a **lifetime measured in years** encompassing **every version** of the application to ever exist, every **hardware failure** and every **human mistake** ever made.
- A data system answers **questions based on information** that was **acquired in the past**

■ How do we build (data) systems today – Today's Architectures





Lack of Human Fault Tolerance

Bugs will be deployed to production over the lifetime of a data system

Operational mistakes will be made

Humans are part of the overall system

- Just like hard disks, CPUs, memory, software
- design for human error like you design for any other fault

Examples of human error

- Deploy a bug that increments counters by two instead of by one
- Accidentally delete data from database
- Accidental DOS on important internal service

Worst two consequences: **data loss** or **data corruption**

As long as an error **doesn't lose or corrupt** good data, you can **fix** what went wrong

■ Lack of Human Fault Tolerance – Immutability vs. Mutability

An immutable system captures historical records of events

Each **event** happens at a **particular time** and is **always true**

Name	City	Timestamp
Guido	Berne	1.8.1999
Albert	Zurich	10.5.1988



Name	City	Timestamp
Guido	Berne	1.8.1999
Albert	Zurich	10.5.1988
Guido	Basel	1.4.2013

The **U** and **D** in CRUD

A mutable system updates the current state of the world

Mutable systems inherently **lack** human fault-tolerance

Name	City
Guido	Berne
Albert	Zurich



Name	City
Guido	Basel
Albert	Zurich

Immutability restricts the range of errors causing data loss/data

corruption Vastly more **human fault-tolerant**

Conclusion: Your **source of truth** should always be


immutable

2014 © Trivadis

DOAG 2014 | Big Data und Fast Data - Lambda Architektur und deren Umsetzung 19.11.2014

■ A different kind of architecture with immutable source of truth

Instead of using our traditional approach ! why not building data systems like this



■ How to create the views on the Immutable data?

On the fly




Materialized, i.e. Pre-computed





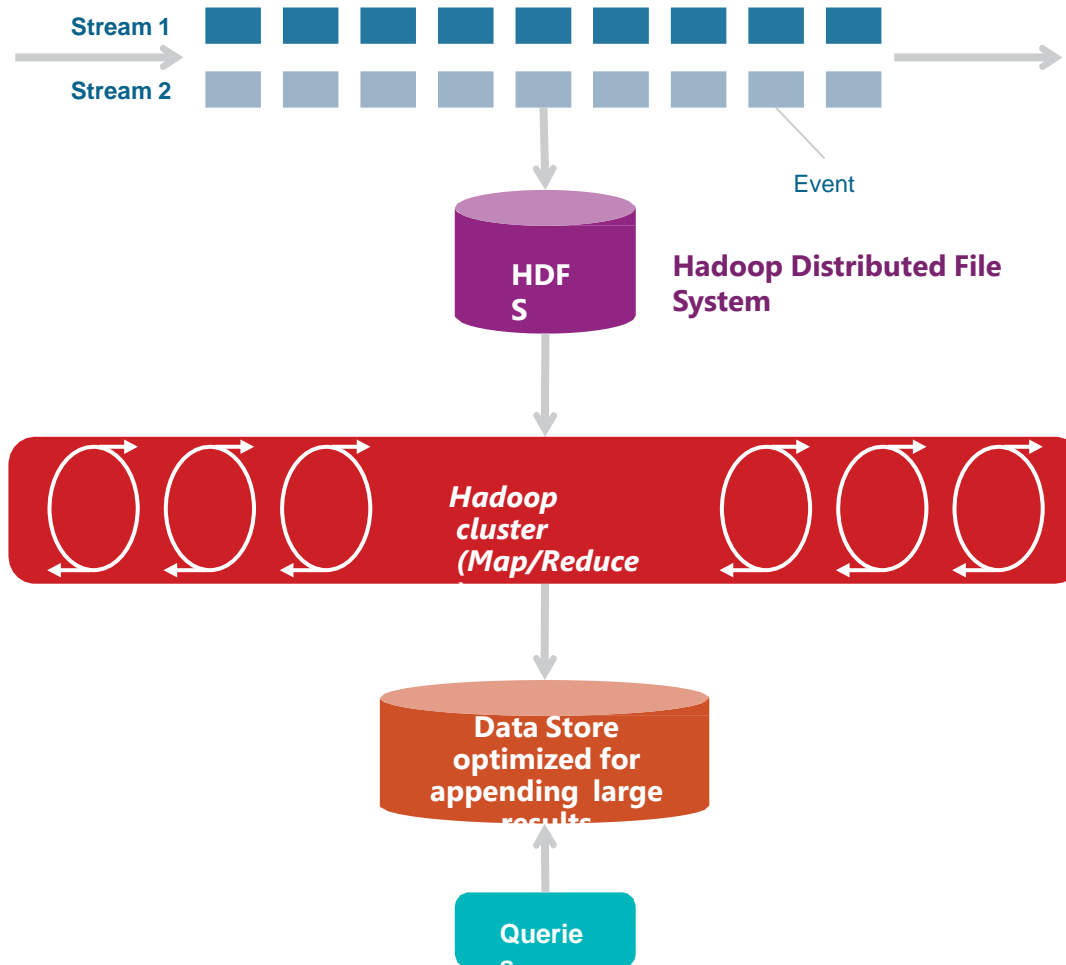
(Big) Data Processing



How to compute the materialized views ?

How to compute queries from the views ?

■ Today Big Data Processing means Batch Processing !





Big Data Processing - Batch

Favorite Product List

1.2.13	Add	iPAD 64GB
10.3.13	Add	Sony RX-100
11.3.13	Add	Canon GX-10
11.3.13	Remove	Sony RX-100
12.3.13	Add	Nikon S-100
14.4.13	Add	BoseQC-15
15.4.13	Add	MacBook Pro
		15
20.4.13	Remove	Canon GX10

**Raw information =>
data**



Current Favorite Product List

iPAD 64GB
Nikon S-100
BoseQC-15
MacBook Pro
15



Current Product Count

4

**Information =>
derived**

■ Big Data Processing – Batch

Favorite Product List Changes

1.2.13	Add	iPAD 64GB
10.3.13	Add	Sony RX-100
11.3.13	Add	Canon GX-10
11.3.13	Remove	Sony RX-100
12.3.13	Add	Nikon S-100
14.4.13	Add	BoseQC-15
15.4.13	Add	MacBook Pro 15
20.4.13	Remove	Canon GX10
Now	Add	Canon Scanner

derive)

Current Favorite Product List

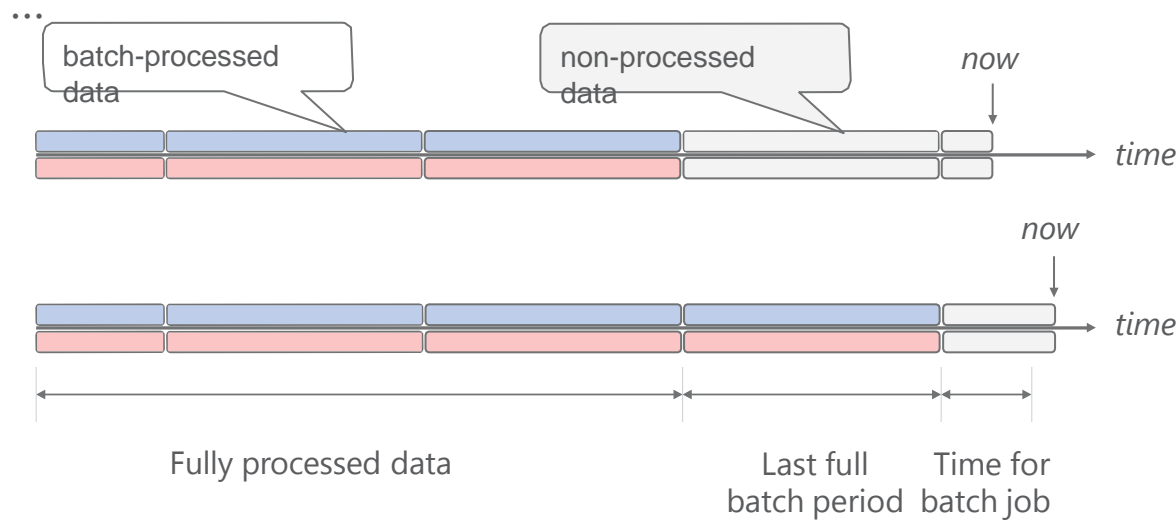
iPAD 64GB
Nikon S-100
BoseQC-15
MacBook Pro 15

derive)

Current Product Count

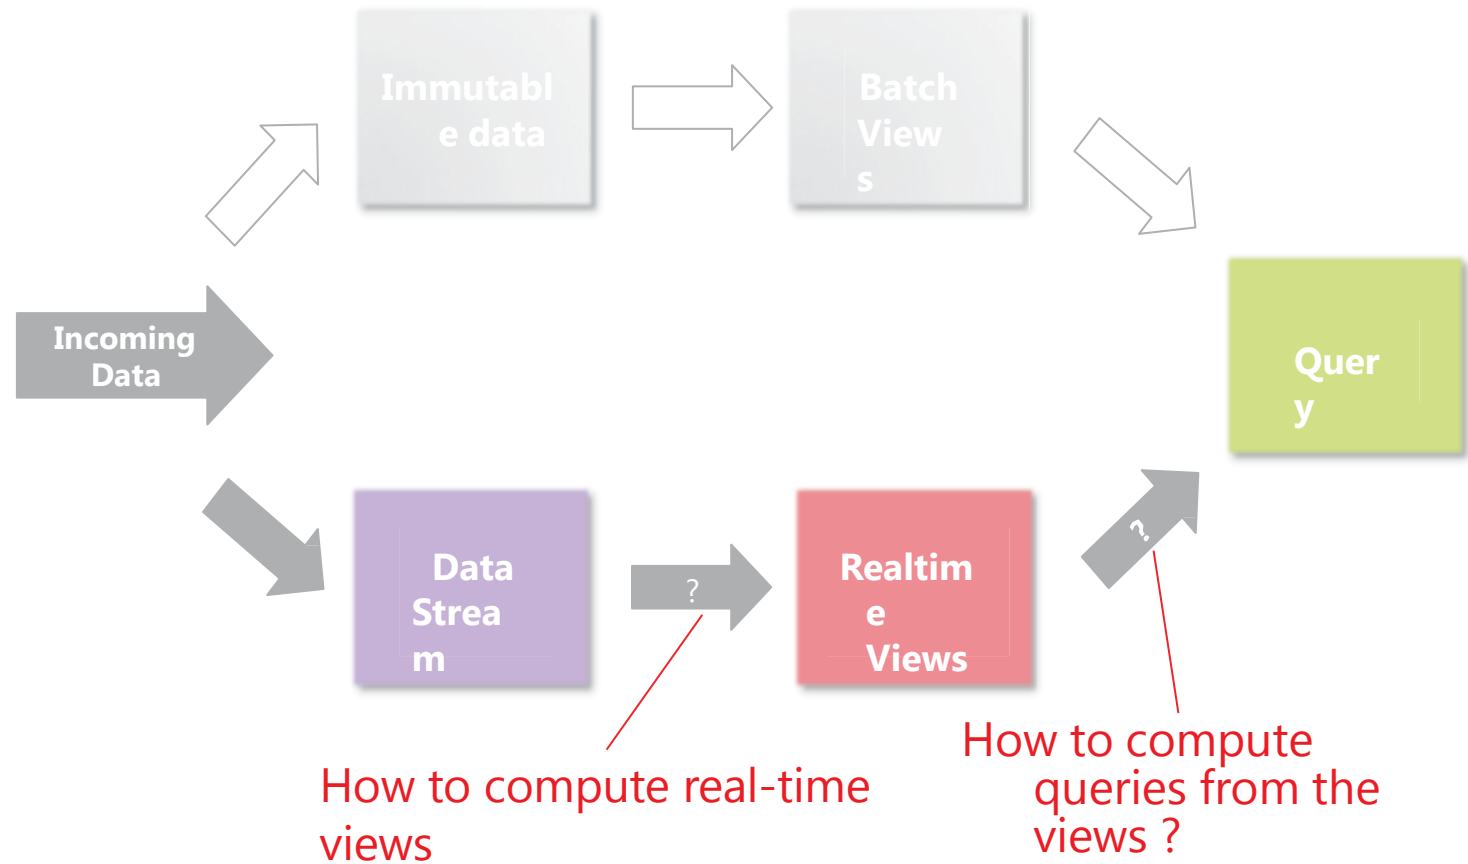
4

But we are not done yet



- Using only batch processing, leaves you always with a portion of non-processed data.

Big Data Processing - Adding Real-Time



Big Data Processing - Adding Real-T

Immutable

View

Data

Quer

Favorite Product List

Changes


1.2.13	Add	iPAD 64GB
10.3.13	Add	Sony RX-100
11..3.13	Add	Canon GX-10
11.3.13	Remove	Sony RX-100
12.3.13	Add	Nikon S-100
14.4.13	Add	BoseQC-15
15.4.13	Add	MacBook Pro 15
20.4.13	Remove	Canon GX10
Now	Add	Canon Scanner

Current Favorite Product List

iPAD 64GB
Nikon S-100
BoseQC-15
MacBook Pro 15

Current Product Count

5



Stream of Favorite Product List Changes

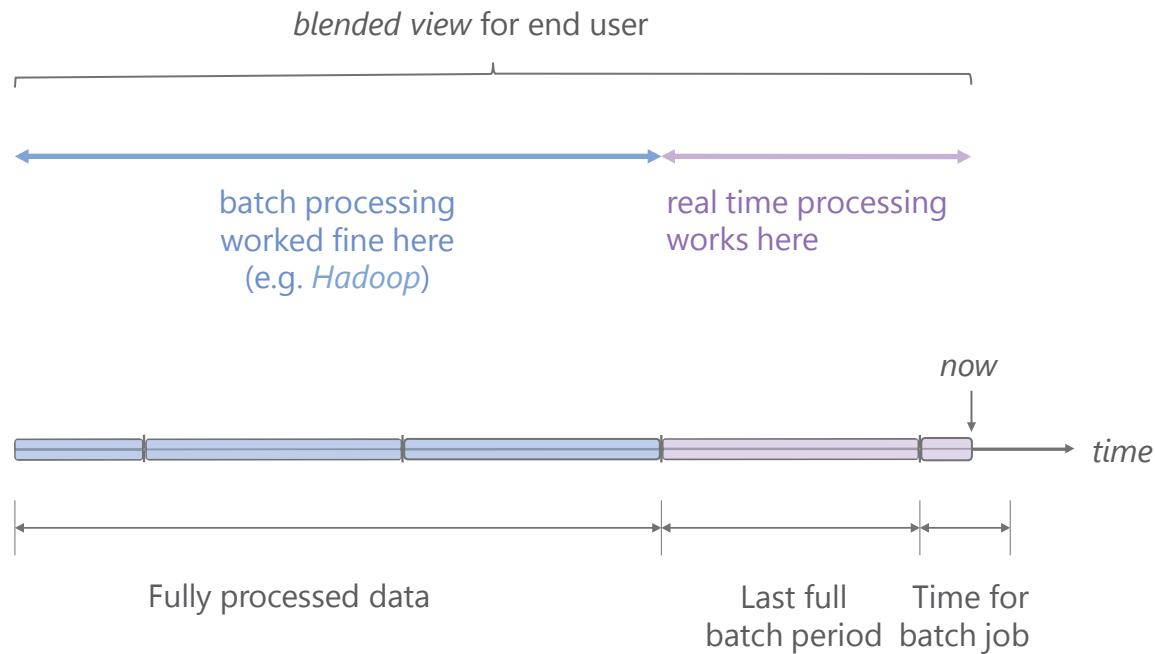

Add Canon Scanner

compute

Now Canon Scanner

query

■ Big Data Processing - Batch & Real Time



Adapted from Ted Dunning (March 2012):
<http://www.youtube.com/watch?v=7PcmbI5aC20>

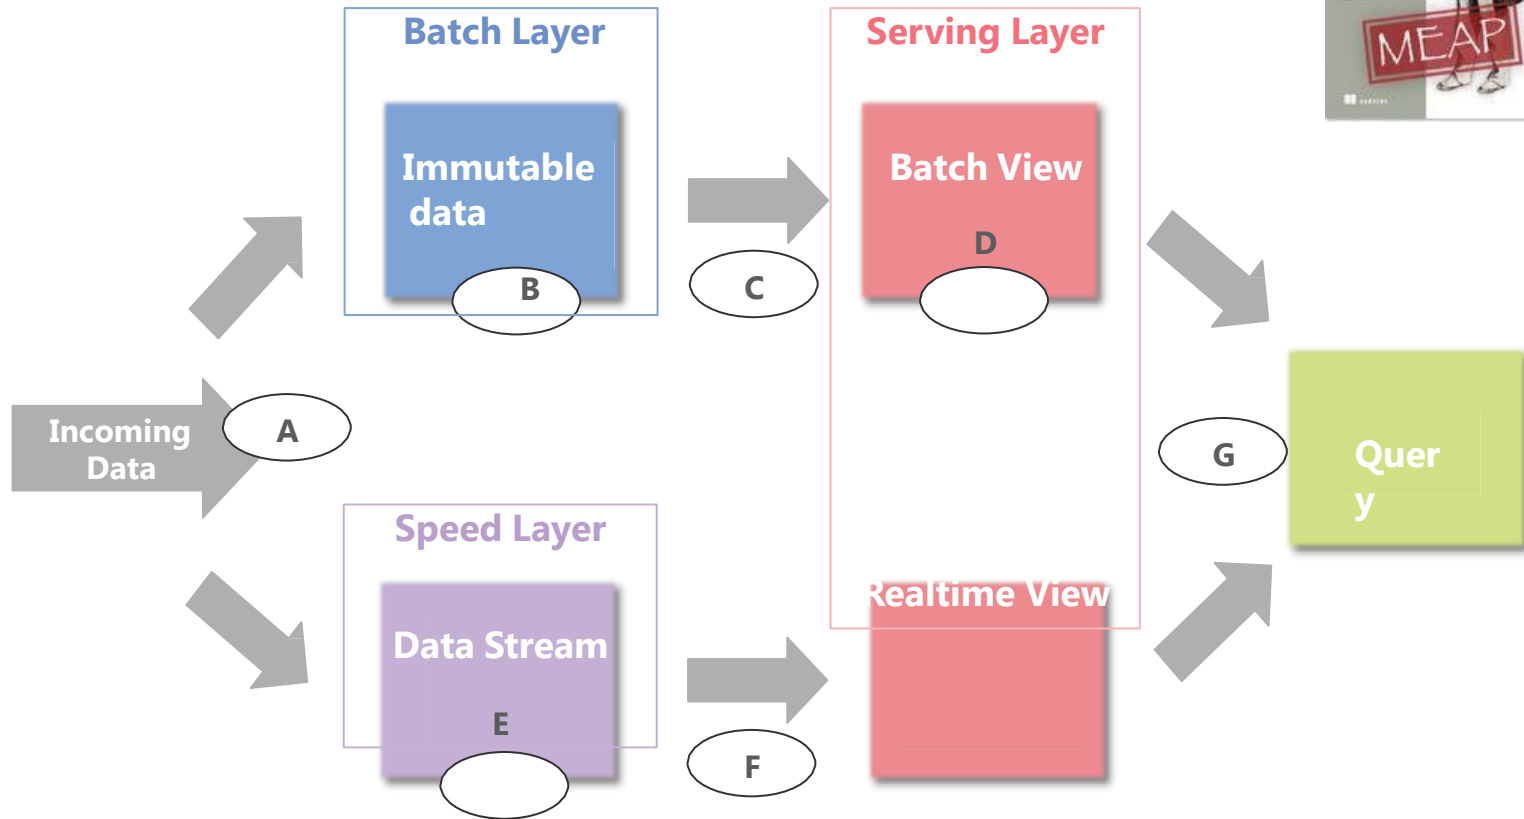



AGENDA

1. Big Data and Fast Data, what is it?
2. Architecting (Big) Data Systems
- 3. The Lambda Architecture**
4. The Use Case and the Implementation
5. Summary and Outlook

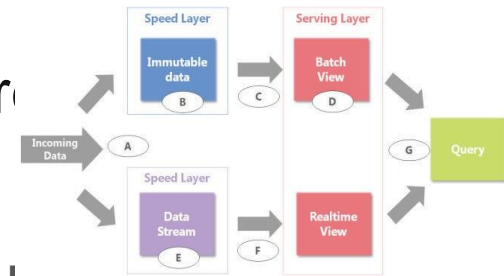


Lambda Architecture



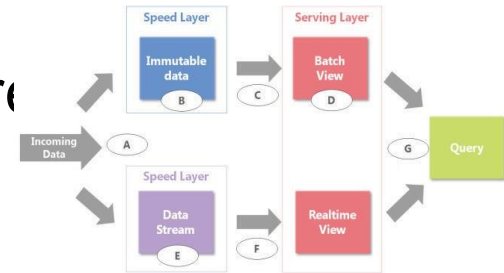
Lambda => Query = function(all data)

Lambda Architecture



- A. All data is sent to **both** the **batch** and **speed layer**
- B. Master data set is an **immutable, append-only** set of data
- C. Batch layer **pre-computes** query functions from scratch, result is called Batch Views. Batch layer **constantly re-computes** the batch views.
- D. Batch views are **indexed** and **stored** in a **scalable database** to get particular values very quickly. Swaps in new batch views when they are available
- E. Speed layer **compensates** for the high latency of updates to the Batch Views
- F. Uses fast **incremental algorithms** and read/write databases to produce real- time views
- G. Queries are resolved by getting results from **both** batch and real-time views

Lambda Architecture



Batch

Stores the immutable constantly growing dataset
Computes arbitrary views from this dataset using
BigData technologies (can take hours)
Can be always recreated

Speed

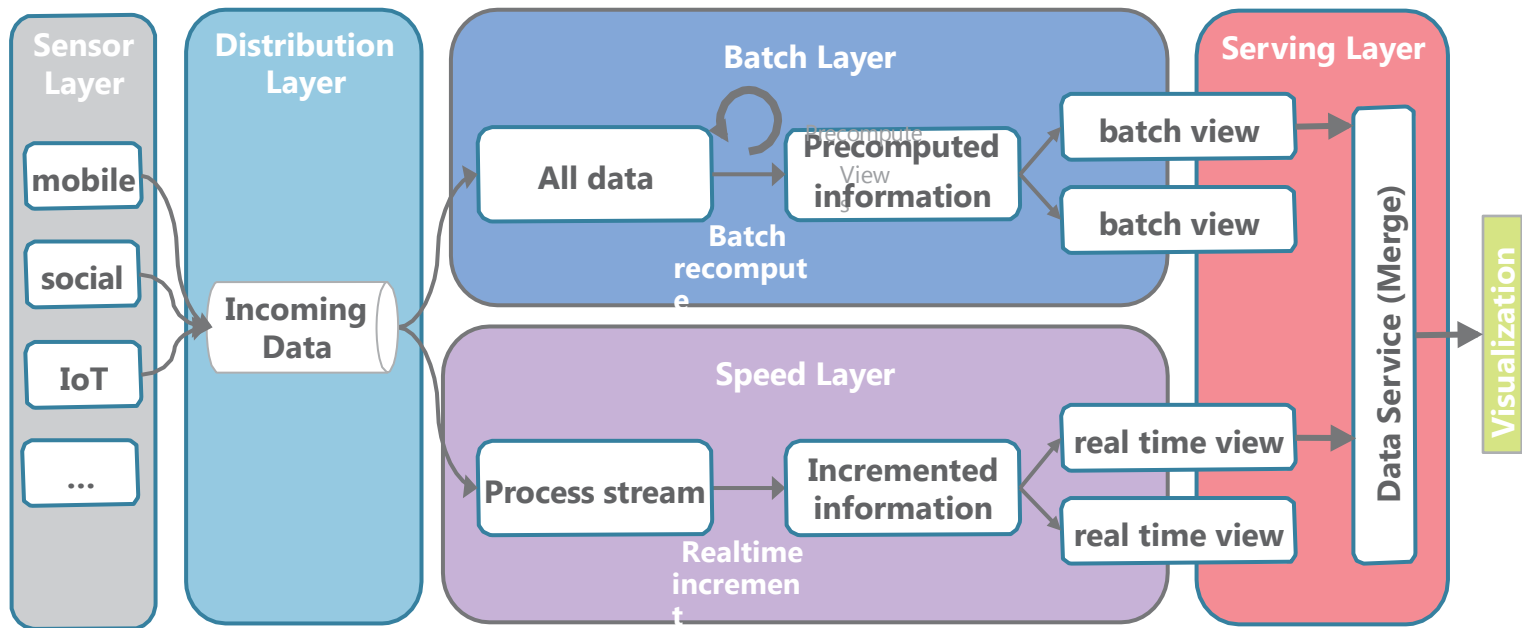
Computes the views from the constant stream of data it
receives Needed to compensate for the high latency of
the batch layer Incremental model and views are

Serving

transient
Responsible for indexing and exposing the pre-
computed batch views so that they can be queried
Exposes the incremented real-time views
Merges the batch and the real-time views into a
consistent result



Lambda Architecture



Adapted from: Marz, N. & Warren, J. (2013) Big Data. Manning.



AGENDA

1. Big Data and Fast Data, what is it?
2. Architecting (Big) Data Systems
3. The Lambda Architecture
- 4. Use Case and the Implementation**
5. Summary and Outlook



Project Definition

- Build a platform for analyzing Twitter communications in retrospective and in real-time
- Scalability and ability for future data fusion with other information is a must
- Provide a Web-based access to the analytical information
- Invest into new, innovative and not widely-proven technology
 - PoC environment, a pre-invest for future systems

Anatomy of a tweet

Time

Space

Content

Social

Technic

```
{
  "created_at": "Sun Aug 18 14:29:11 +0000 2013",
  "id": "369103686938546176", "id_str": "369103686938546176",
  "text": "Baloncesto preparaci\u00f3n Eslovenia, Rajoy derrota a Merkel. #quelosepash",
  "source": "\u003ca href=\\"http://twitter.com/download/iphone\\" rel=\\"nofollow\\" \u003eTwitter for iPhone\u003c/a\u003e",
  "truncated": false,
  "in_reply_to_status_id": null, "in_reply_to_status_id_str": null,
  "in_reply_to_user_id": null, "in_reply_to_user_id_str": null,
  "in_reply_to_screen_name": null, "user": {
    "id": "15032594", "id_str": "15032594",
    "name": "Juan Carlos Romo\u002122",
    "screen_name": "jcsromo", "location": "Sopuerta, Vizcaya", "url": null,
    "description": "Portugalujo, saturado de todo, de baloncesto no. Twitter personal.",
    "protected": false,
    "followers_count": 1331, "friends_count": 1326, "listed_count": 31,
    "created_at": "Fri Jun 06 21:21:22 +0000 2008", "favourites_count": 255,
    "utc_offset": 7200, "time_zone": "Madrid", "geo_enabled": true, "verified": false,
    "statuses_count": 22787, "lang": "es", "contributors_enabled": false,
    "is_translator": false,
    ...
    "profile_image_url_https": "https://si0.twimg.com/profile_images/2649762203/be4973d9eb457a45077897879c47c8b7_normal.jpeg",
  },
}
```

```
  "profile_banner_url": "https://pbs.twimg.com/profile_banners/15032594/1371570460",
  "profile_link_color": "2FC2EF",
  "profile_sidebar_border_color": "FFFFFF", "profile_sidebar_fill_color": "252429",
  "profile_text_color": "666666", "profile_use_background_image": true,
  "default_profile": false, "default_profile_image": false, "following": null,
  "follow_request_sent": null, "notifications": null,

  "geo": {
    "type": "Point", "coordinates": [43.28261499, -2.96464655],
    "coordinates": { "type": "Point", "coordinates": [-2.96464655, 43.28261499] },
    "place": { "id": "cd43ea85d651af92",
    "url": "https://api.twitter.com/1.1/geo/vid/cd43ea85d651af92.json",
    "place_type": "city",
    "name": "Bilbao", "full_name": "Bilbao, Vizcaya", "country_code": "ES",
    "country": "Espa\u00f1a",
    "bounding_box": { "type": "Polygon", "coordinates": [[ [-2.9860102, 43.2136542],
    [-2.9860102, 43.2901452], [-2.8803248, 43.2901452], [-2.8803248, 43.2136542]] ] },
    "attributes": {} },
  },

  "contributors": null, "retweet_count": 0, "favorite_count": 0,

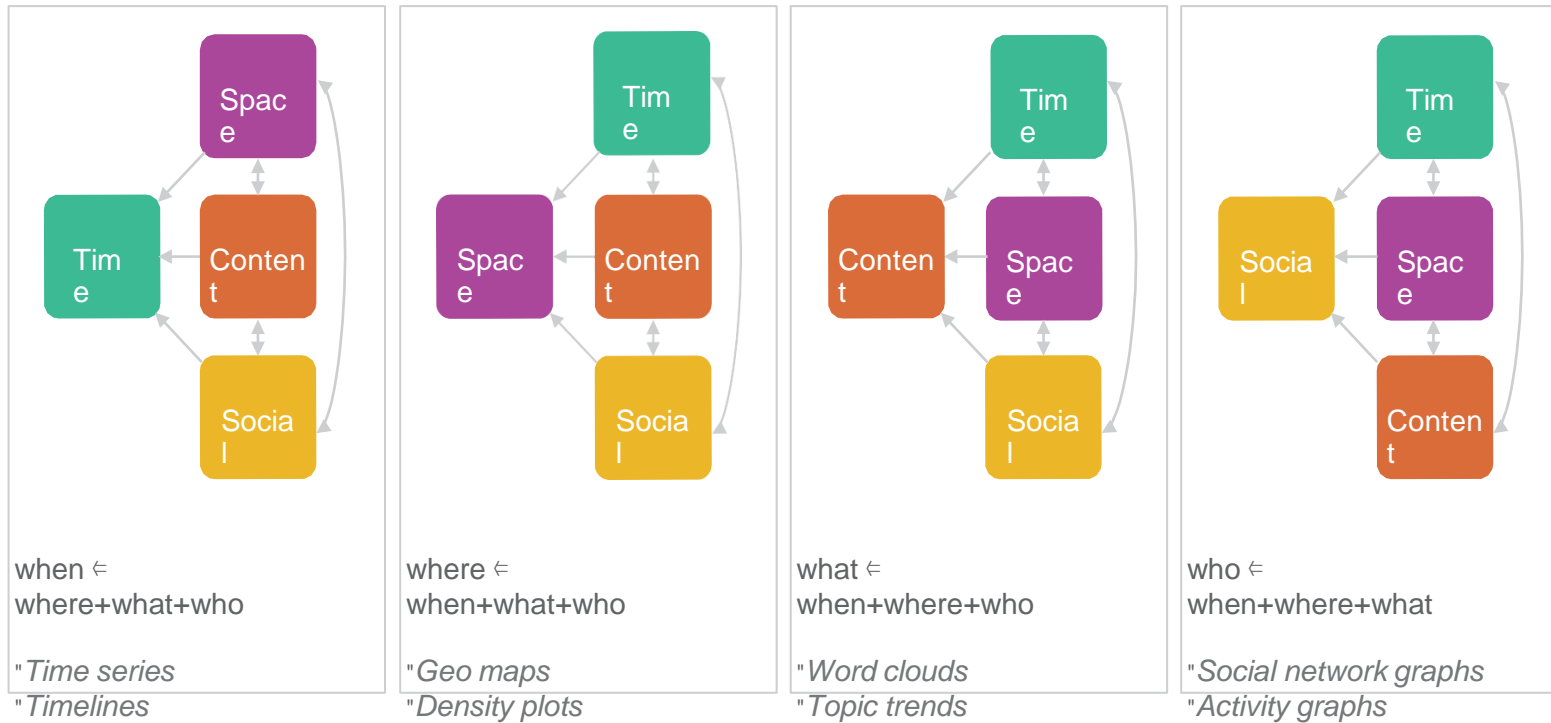
  "entities": { "hashtags": [ { "text": "quelosepash", "indices": [58, 70] } ], "symbols": [],
  "urls": [],

  "user_mentions": [], "favorited": false, "retweeted": false, "filter_level": "medium",

  "lang": "es"
}
```



Views on Tweets in four dimensions



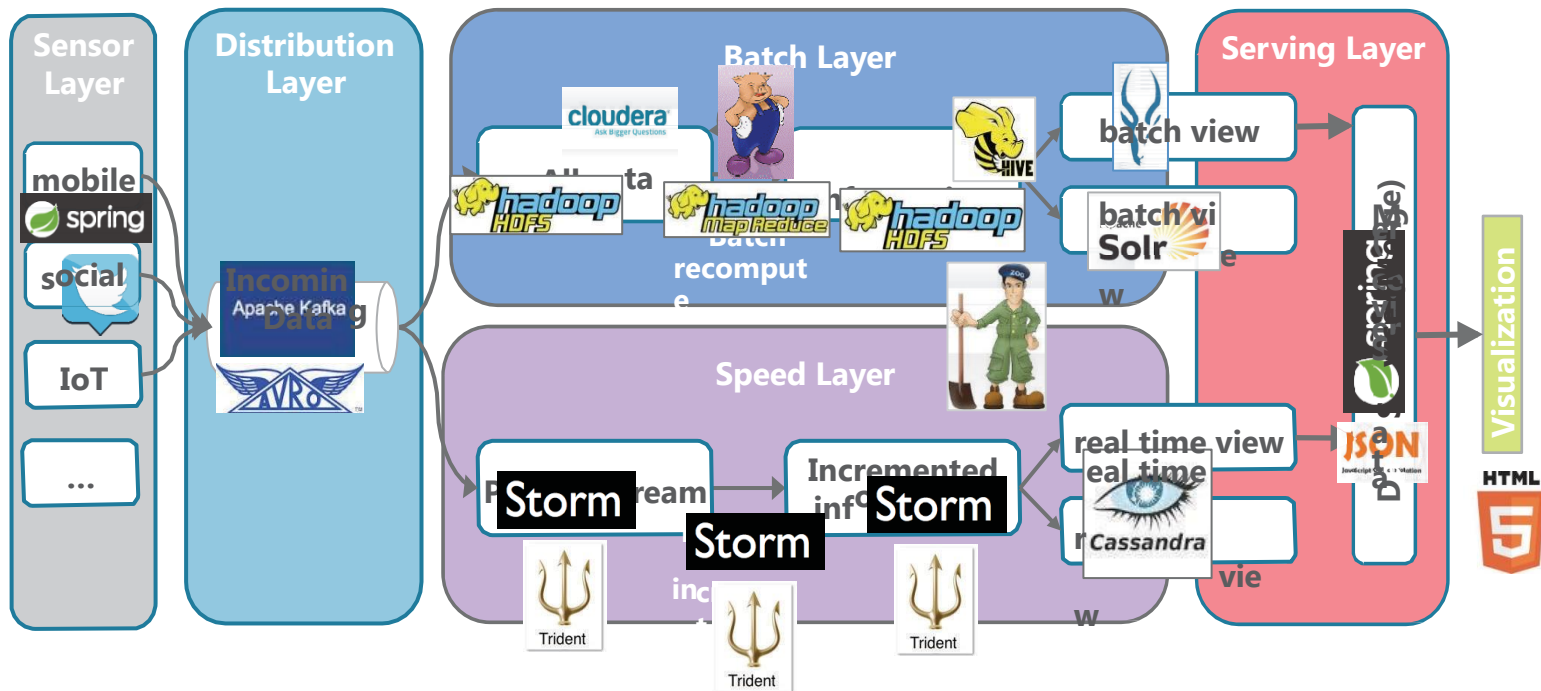


Accessing Twitter

Quelle	Limitierungen	Zugang
Twitter's Search API	3200 / user 5000 / keyword 180 Anfragen / 15 Minuten	gratis
Twitter's Streaming API	1%-40% des Volumens	gratis
DataSift	keine	0.15 -0.20\$ / unit
Gnip	keine	Auf Anfrage

Lambda Architecture

Open Source Frameworks for implementing a Lambda Architecture





Lambda Architecture in Action

Twitter Horsebird Client (hbc)

Distribution

- Twitter Java API over Streaming API
-

Popular Java Framework used to modularize part of the logic (sensor and serving layer)

Apache Kafka

- Simple messaging framework based on file system to distribute information to both batch and speed layer

Apache Avro

- Serialization system for efficient cross-language RPC and persistent data storage

JSON

- open standard format that uses human-readable text to transmit data objects consisting of attribute–value pairs.

Cloudera

Distribution of Apache Hadoop: HDFS, MapReduce, Hive, Flume, Pig, Impala

Cloudera

Impala

- distributed query execution engine that runs against data stored in HDFS and HBase

Apache Zookeeper

- Distributed, highly available coordination service. Provides primitives such as distributed locks

Apache Storm &

Trident

- distributed, fault-tolerant realtime computation system

Apache

Cassandra

- distributed database management system designed to handle large amounts of data across many commodity servers, providing high availability with no single point of failure



Facts & Figures

14 active twitter feeds

~ 14 million tweets/day (> 5 billion tweets/year)

~ 8 GB/day raw data, compressed (2 DVDs)

66 GB storage capacity / day (replication & views/results included)

Cluster of 10 nodes

- ~100 processors

- ~40 TB HD capacity in total; 46% used

- >500 GB RAM

Currently in total

- 2.7 TB Raw Data

- 1.1 TB Pre-Processed data in Impala

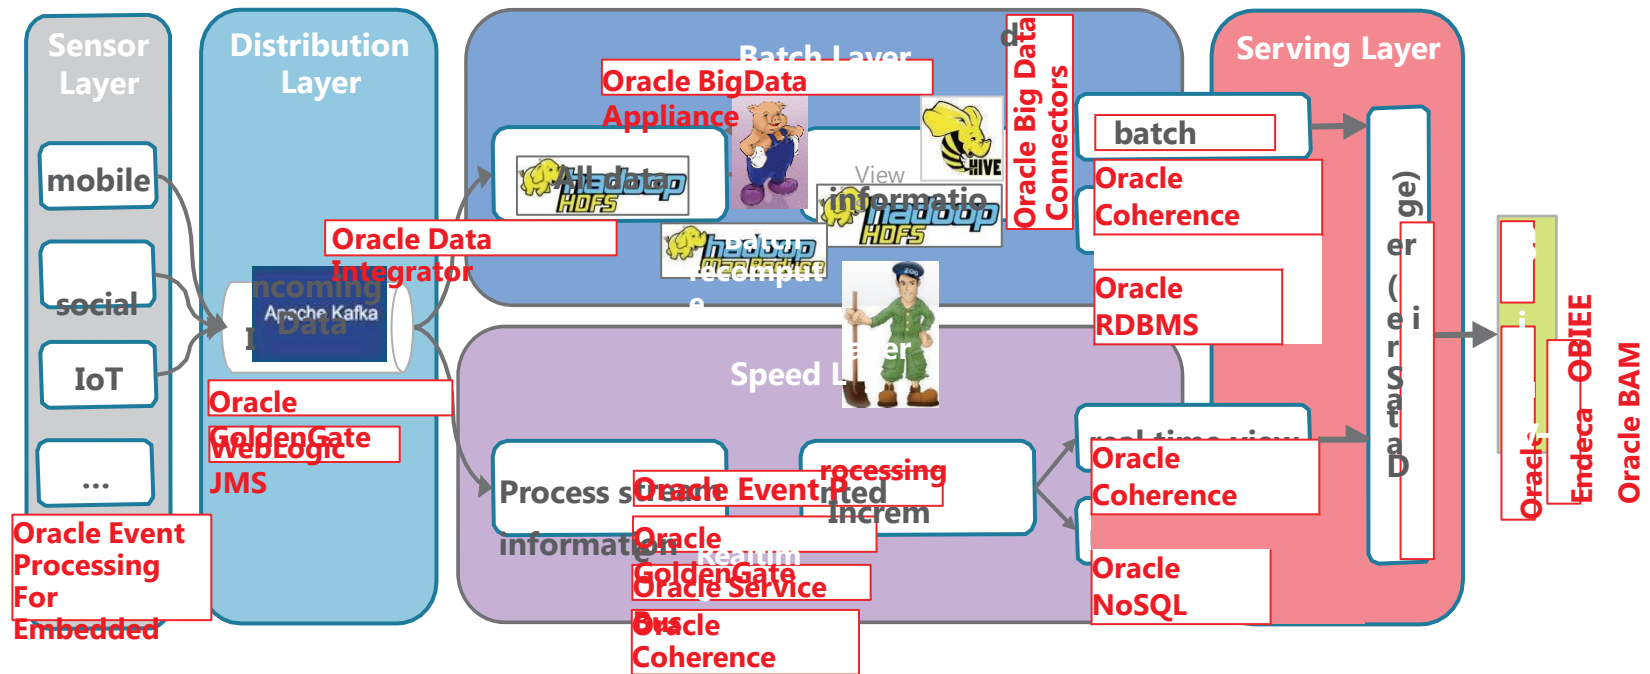
- 1 TB Solr indices for full text search

Cloudera 4.7.0 with Hadoop, Pig, Hive, Impala and Solr

Kafka 0.7, Storm 0.9, DataStax Enterprise Edition

Lambda Architecture with Oracle Product Stack

Possible implementation with Oracle Product stack





AGENDA

1. Big Data and Fast Data, what is it?
2. Architecting (Big) Data Systems
3. The Lambda Architecture
4. Use Case and the Implementation
- 5. Summary and Outlook**

Summary – The lambda architecture





- Can discard batch views and real-time views and recreate everything from scratch
- Mistakes corrected via re-computation
- Scalability through platform and distribution
- Data storage layer optimized independently from query resolution layer
- Still in a early stage !. But a very interesting idea!
 - Today a zoo of technologies are needed => Infrastructure group might not like it
 - Better with so-called Hadoop distributions and Hadoop V2 (YARN)



Alternative Approaches – Motivation

Data Sharing in Map Reduce

!






Alternative Approaches – Motivation

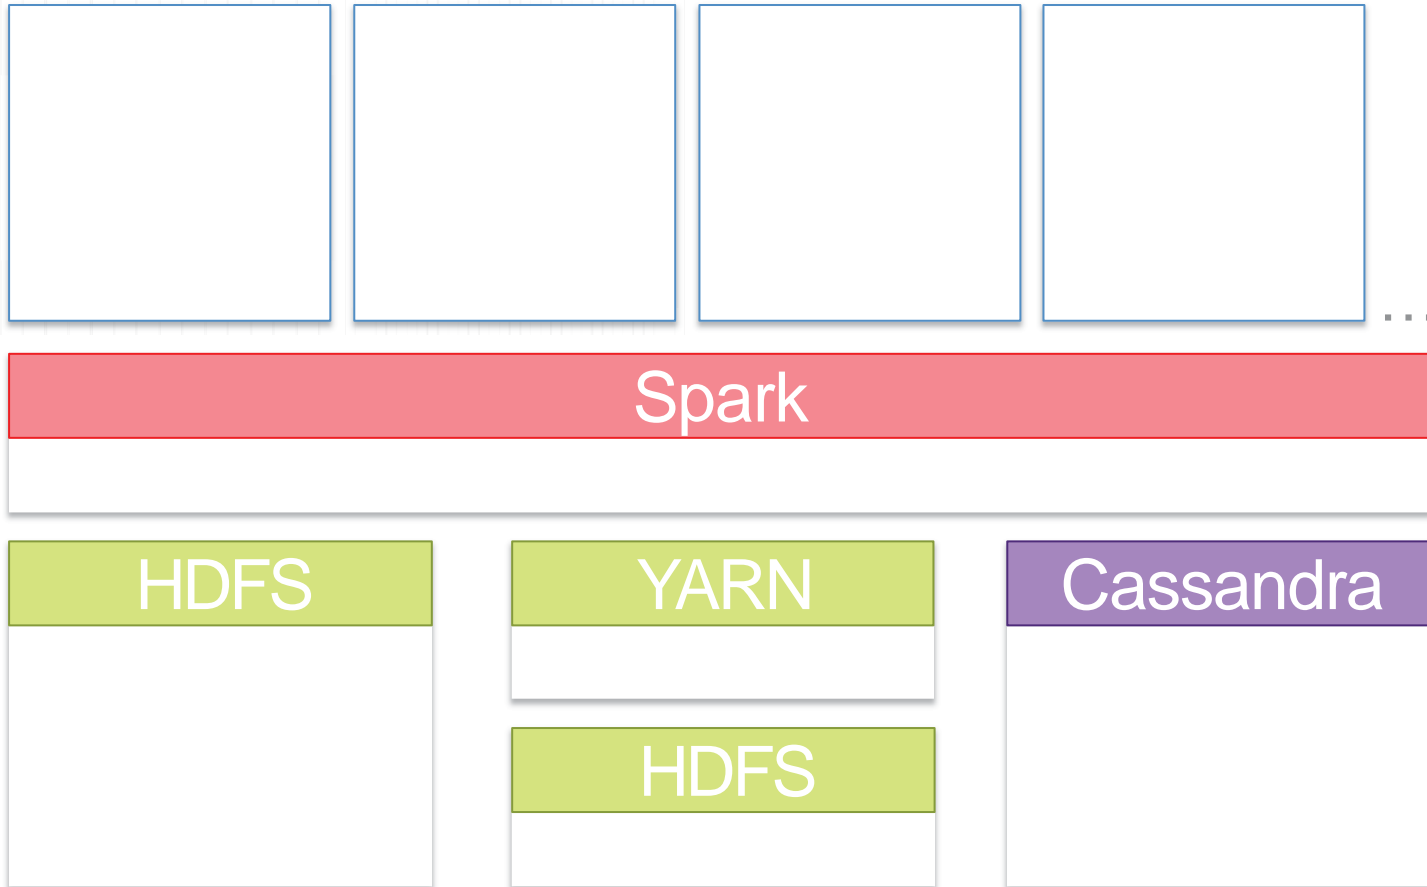
What we would like

!

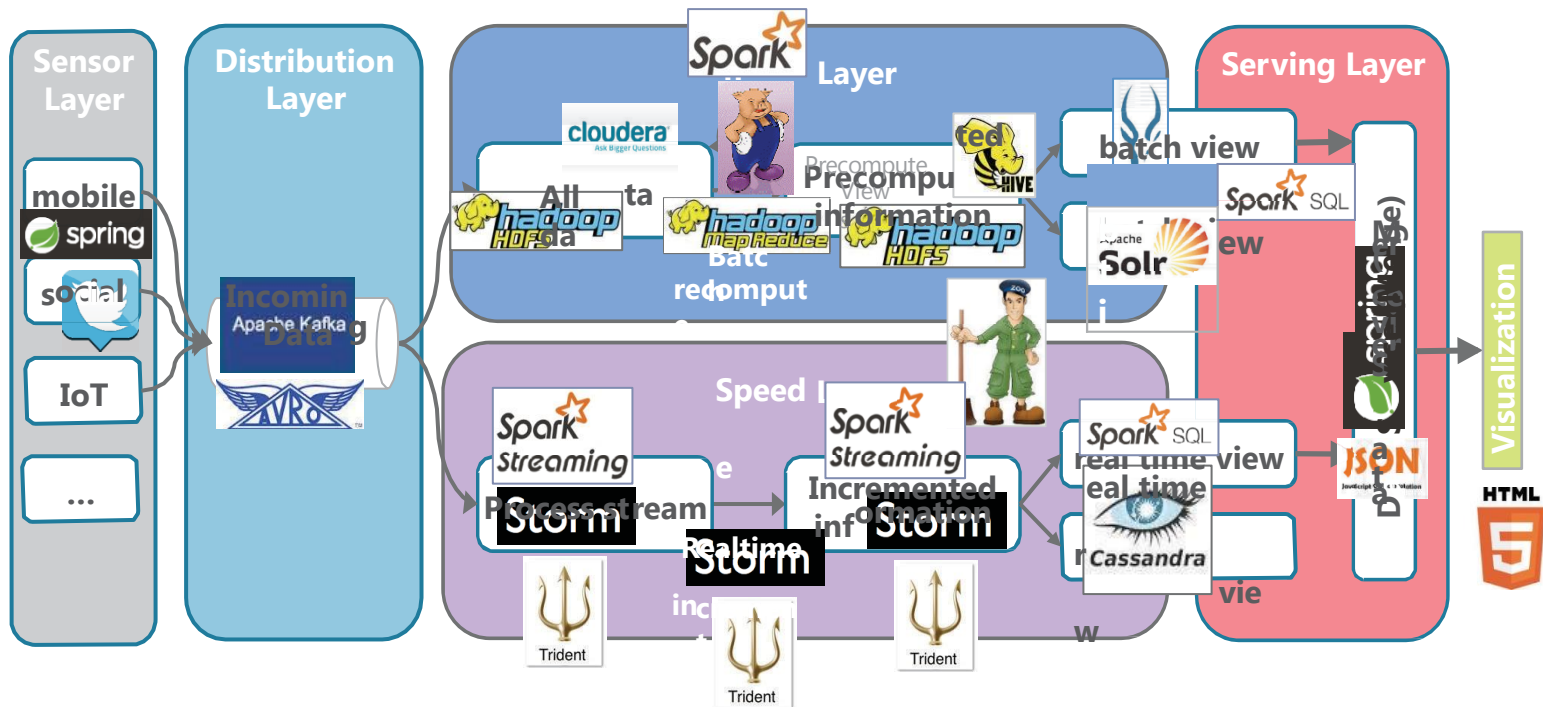




Alternatives – Apache Spark

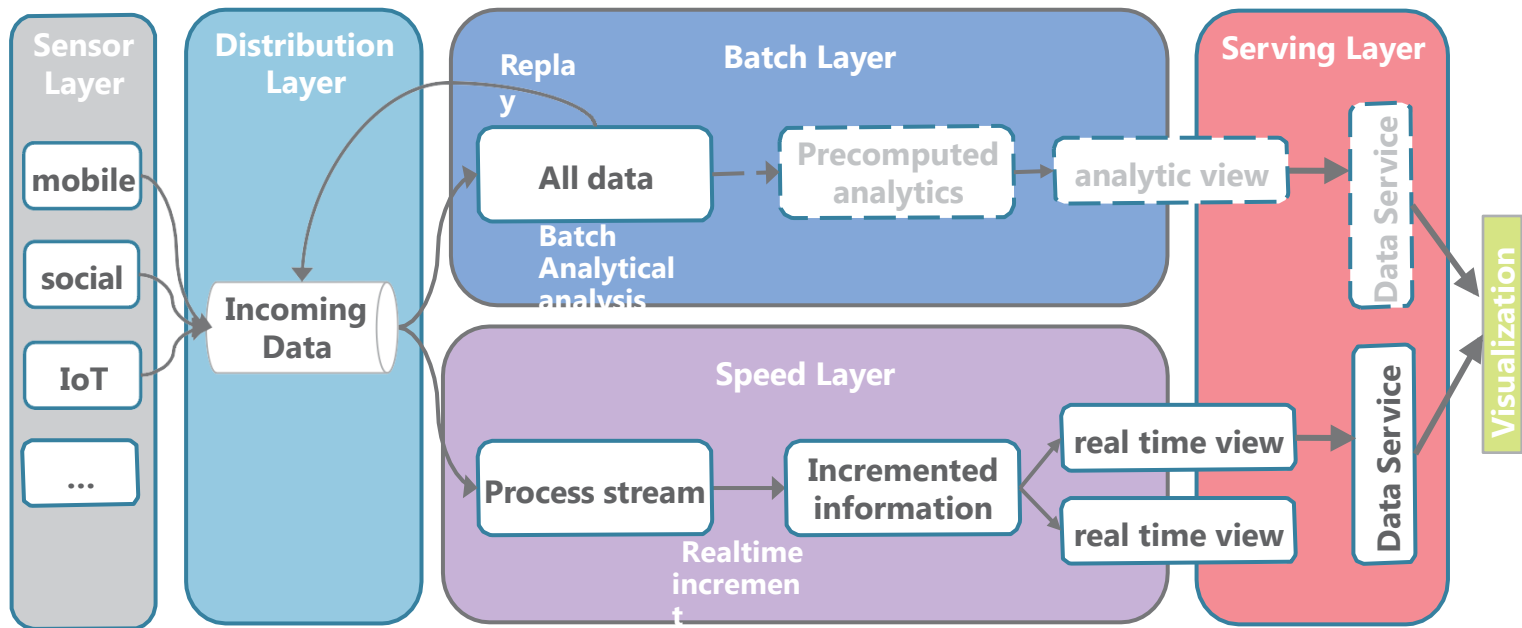


Alternative Technologies – Apache Spark





“Kappa Architecture”



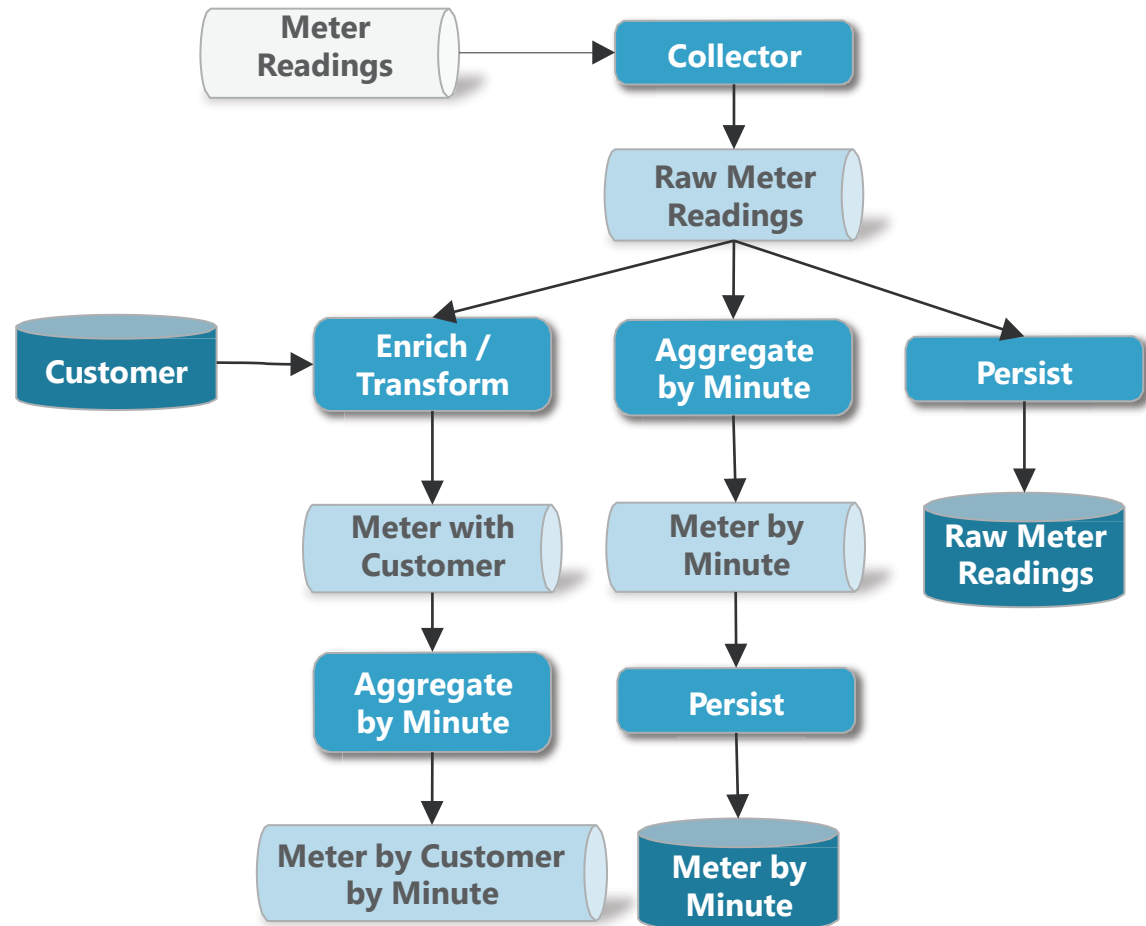
Adapted from: Marz, N. & Warren, J. (2013) Big Data. Manning.

Unified Log Processing Architecture

Stream processing allows for computing feeds off of other feeds

Derived feeds are no different than original feeds they are computed off

Single deployment of “Unified Log” but logically different feeds





Thank You :)