

CSE 486/586 Distributed Systems

Paxos --- 1

Steve Ko
Computer Sciences and Engineering
University at Buffalo

CSE 486/586, Spring 2014

Recap

- NFS
 - Caching with write-through policy at close()
 - Stateless server
- One power efficient design: FAWN (Fast Array of Wimpy Nodes)
 - Embedded CPUs & Flash storage
 - Write problem: block erasure first
 - FTL presents a logical structure different from the physical structure. Physically, it's log-structured.

FTL: Flash Translation Layer

CSE 486/586, Spring 2014

2

Paxos

- A consensus algorithm
 - Known as one of the most efficient & elegant consensus algorithms
 - If you stay close to the field of distributed systems, you'll hear about this algorithm over and over.
- What? Consensus? What about FLP (the impossibility of consensus)?
 - Obviously, it doesn't solve FLP.
 - It relies on failure detectors to get around it.
- Plan
 - Brief history (with a lot of quotes)
 - The protocol itself
 - How to "discover" the protocol

(FLP: Fischer, Lynch, and Paterson)

CSE 486/586, Spring 2014

3

Brief History

- Developed by Leslie Lamport (from the Lamport clock)
- "A fault-tolerant file system called Echo was built at SRC in the late 80s. The builders claimed that it would maintain consistency despite any number of non-Byzantine faults, and would make progress if any majority of the processors were working."
- "I decided that what they were trying to do was impossible, and set out to prove it. Instead, I discovered the Paxos algorithm."
- "I decided to cast the algorithm in terms of a parliament on an ancient Greek island (Paxos)."

CSE 486/586, Spring 2014

4

Brief History

- The paper abstract:
 - "Recent archaeological discoveries on the island of Paxos reveal that the parliament functioned despite the peripatetic propensity of its part-time legislators. The legislators maintained consistent copies of the parliamentary record, despite their frequent forays from the chamber and the forgetfulness of their messengers. The Paxos parliament's protocol provides a new way of implementing the state-machine approach to the design of distributed systems."
- "I gave a few lectures in the persona of an Indiana-Jones-style archaeologist."
- "My attempt at inserting some humor into the subject was a dismal failure. People who attended my lecture remembered Indiana Jones, but not the algorithm."

CSE 486/586, Spring 2014

5

Brief History

- People thought that Paxos was a joke.
- Lamport finally published the paper 8 years later in 1998 after it was written in 1990.
 - Title: "The Part-Time Parliament"
- People did not understand the paper.
- Lamport gave up and wrote another paper that explains Paxos in simple English.
 - Title: "Paxos Made Simple"
 - Abstract: "The Paxos algorithm, when presented in plain English, is very simple."
- Still, it's not the easiest algorithm to understand.
- So people started to write papers and lecture notes to explain "Paxos Made Simple." (e.g., "Paxos Made Moderately Complex", "Paxos Made Practical", etc.)

CSE 486/586, Spring 2014

6

Review: Consensus

- How do people agree on something?
 - Q: should Steve give an A to everybody taking CSE 486/586?
 - Input: everyone says either yes/no.
 - Output: an agreement of yes or no.
 - FLP: this is impossible even with one-faulty process and arbitrary delays.
- Many distributed systems problems can cast into a consensus problem
 - Mutual exclusion, leader election, total ordering, etc.
- Paxos
 - How do multiple processes agree on a value?
 - Under failures, network partitions, message delays, etc.

CSE 486/586, Spring 2014

7

Review: Consensus

- People care about this!
- Real systems implement Paxos
 - Google Chubby
 - MS Bing cluster management
 - Etc.
- Amazon CTO Werner Vogels (in his blog post “Job Openings in My Group”)
 - “What kind of things am I looking for in you?”
 - “You know your distributed systems theory: You know about logical time, snapshots, stability, message ordering, but also acid and multi-level transactions. You have heard about the FLP impossibility argument. You know why failure detectors can solve it (but you do not have to remember which one diamond-w was). You have at least once tried to understand Paxos by reading the original paper.”

CSE 486/586, Spring 2014

8

CSE 486/586 Administrivia

- PA3 scores will be posted by tonight.
- Midterm scores will be posted by tonight.
- PA4 released.
 - Tester will be released soon.
 - A small correction will be posted as well.

CSE 486/586, Spring 2014

9

Paxos Assumptions & Goals

- The network is *asynchronous* with message delays.
- The network can *lose or duplicate* messages, but *cannot corrupt* them.
- Processes can *crash*.
- Processes are *non-Byzantine* (only crash-stop).
- Processes have *permanent storage*.
- Processes can *propose* values.
- The goal: every process agrees on a value out of the proposed values.

CSE 486/586, Spring 2014

10

Desired Properties

- Safety
 - Only a value that has been proposed can be chosen
 - Only a single value is chosen
 - A process never learns that a value has been chosen unless it has been
- Liveness
 - Some proposed value is eventually chosen
 - If a value is chosen, a process eventually learns it

CSE 486/586, Spring 2014

11

Roles of a Process

- Three roles
- **Proposers**: processes that propose values
- **Acceptors**: processes that accept (i.e., consider) values
 - “Considering a value”: the value is a candidate for consensus.
 - Majority acceptance → choosing the value
- **Learners**: processes that learn the outcome (i.e., chosen value)
- In reality, a process can be any one, two, or all three.
- Important requirements
 - The protocol should work under process failures and with delayed and lost messages.
 - The consensus is reached via a majority ($> \frac{1}{2}$).

CSE 486/586, Spring 2014

12

Roles of a Process


- In reality, a process can be any one, two, or all three.
- Important requirements
 - The protocol should work under process failures and with delayed and lost messages.
 - The consensus is reached via a majority ($> \frac{1}{2}$).
- Example: a replicated state machine
 - All replicas agree on the order of execution for concurrent transactions
 - All replicas assume all roles, i.e., they can each propose, accept, and learn.

CSE 486/586, Spring 2014

13

First Attempt

- Let's just have one acceptor, choose the first one that arrives, & tell the proposers about the outcome.



- What's wrong?


- Single point of failure!

CSE 486/586, Spring 2014

14

Second Attempt

- Let's have multiple acceptors; each accepts the first one; then all choose the majority and tell the proposers about the outcome.




- What's wrong? (next slide)

CSE 486/586, Spring 2014

15

Second Attempt

- One example, but many other possibilities



CSE 486/586, Spring 2014

16

Paxos

- Let's have each acceptor accept (i.e., consider) **multiple proposals**.
 - An acceptor accepting a proposal doesn't mean it will be chosen. A majority should accept it.
 - "Hope" that one of the multiple accepted proposals will have a vote from a majority (will get back to this later)
- Paxos: how do we select one value when there are multiple acceptors accepting multiple proposals?

CSE 486/586, Spring 2014

17

Paxos Protocol Overview

- A proposal should have an ID.
 - (proposal #, value) == (N, V)
 - The proposal # strictly increasing and globally unique across all proposers
- Three phases
 - Prepare phase: a proposer learns previously-accepted proposals from the acceptors.
 - Propose phase: a proposer sends out a proposal.
 - Learn phase: learners learn the outcome.

CSE 486/586, Spring 2014

18

Paxos Protocol Overview

- Rough description of the **proposers**
 - Before a proposer proposes a value, it will ask acceptors if there is any proposed value already.
 - If there is, the proposer will propose the same value, rather than proposing another value.
 - The behavior is **altruistic**: the goal is to reach a consensus, rather than making sure that "my value" is chosen.
- Rough description of the **acceptors**
 - The goal for acceptors is to accept the highest-numbered proposal coming from all proposers.
 - An acceptor tries to accept a value V with the highest proposal number N.
- Rough description of the **learners**
 - All learners are passive and wait for the outcome.

CSE 486/586, Spring 2014

19

Paxos Phase 1

- A proposer chooses its proposal number N and sends a **prepare request** to acceptors.
 - "Hey, have you accepted any proposal yet?"
- An acceptor needs to reply:
 - If it accepted anything, the accepted proposal and its value with the **highest proposal number less than N**
 - A **promise to not accept** any proposal numbered **less than N** any more (to make sure that it doesn't alter the result of the reply).

CSE 486/586, Spring 2014

20

Paxos Phase 2

- If a proposer receives a reply from a majority, it sends an **accept request** with the proposal (N, V).
 - V: the value from the **highest proposal number N** from the replies (i.e., the accepted proposals returned from acceptors in phase 1)
 - Or, if no accepted proposal was returned in phase 1, a new value to propose.
- Upon receiving (N, V), acceptors either:
 - **Accept** it
 - Or, **reject** it if there was another prepare request with N' higher than N, and it replied to it.

CSE 486/586, Spring 2014

21

Paxos Phase 3

- Learners need to know which value has been chosen.
- Many possibilities
- One way: have each acceptor respond to all learners
 - Might be effective, but expensive
- Another way: elect a "distinguished learner"
 - Acceptors respond with their acceptances to this process
 - This distinguished learner informs other learners.
 - Failure-prone
- Mixing the two: a set of distinguished learners

CSE 486/586, Spring 2014

22

Problem: Progress (Liveness)

- **There's a race condition for proposals.**
- P0 completes phase 1 with a proposal number N0
- Before P0 starts phase 2, P1 starts and completes phase 1 with a proposal number N1 > N0.
- P0 performs phase 2, acceptors reject.
- Before P1 starts phase 2, P0 restarts and completes phase 1 with a proposal number N2 > N1.
- P1 performs phase 2, acceptors reject.
- ... (this can go on forever)

CSE 486/586, Spring 2014

23

Providing Liveness

- Solution: **elect a distinguished proposer**
 - I.e., have only one proposer
- If the distinguished proposer can successfully communicate with a majority, the protocol guarantees liveness.
 - I.e., if a process plays all three roles, Paxos can tolerate failures $f < 1/2 * N$.
- Still needs to get around FLP for the leader election, e.g., having a failure detector

CSE 486/586, Spring 2014

24

Summary

- Paxos
 - A consensus algorithm
 - Handles crash-stop failures ($f < 1/2 * N$)
- Three phases
 - Phase 1: prepare request/reply
 - Phase 2: accept request/reply
 - Phase 3: learning of the chosen value

CSE 486/586, Spring 2014

25

Acknowledgements

- These slides contain material developed and copyrighted by Indranil Gupta (UIUC).

CSE 486/586, Spring 2014

26